

Kno.

IT / DSI Technical Datasheet

AI Document Assistant — 100% local & on-premises

100% On-Premises • GDPR by Design • Zero Cloud • Mistral / LLaMA / Phi-3 • Multi-collections • v1.8 • Fdevelopment LTD

Overview

Kno is an AI-powered DMS assistant (Document Management System) running entirely on local infrastructure. It enables natural language queries across your internal document base — contracts, procedures, reports, regulations — without any data leaving your network.

Accessible from any browser on your internal network, Kno installs on a dedicated server and serves all authorized users simultaneously. The AI engine runs entirely locally using open-source GGUF models (Mistral, LLaMA, Phi-3).

- + No data transmitted
- + Local AI model (GGUF)
- + Cited & clickable sources
- + Browser interface
- + Works offline
- + Simultaneous multi-user
- + Multi-collections
- + Encrypted conversations

Recommended Hardware Configuration

Component	Minimum	Recommended	Optimal (responses < 5s)
OS	Windows 10 64-bit or Linux Ubuntu 20+	Windows 11 Pro or Ubuntu 22.04 LTS	Ubuntu Server 22.04 LTS
RAM	8 GB	16 GB	64 GB – 128 GB
CPU	4 cores	8 cores	AMD Ryzen 9 / Intel i9 16+ cores
GPU (AI)	— (CPU only)	NVIDIA RTX 3060 12 GB	NVIDIA RTX 4080/4090 16–24 GB VRAM
Storage	6 GB (+ models)	512 GB NVMe SSD	2 TB NVMe SSD
Network	LAN 100 Mbit	Gigabit Ethernet	Gigabit + HTTPS reverse proxy

Note: CPU-only mode generates ~2–5 tokens/s with Mistral-7B. With RTX 4080/4090: 10–20× faster (3–5 s/response). Apple MacBook Pro M5 Max 64 GB (unified memory) is a silent high-performance alternative.

Compatible AI Models (GGUF format)

Model	Size	RAM	Quality	Recommended use
Mistral-7B-Instruct v0.3 Q4_K_M	4.4 GB	8 GB	★★★★★	Reference — general use, French language
LLaMA-3.1-8B-Instruct Q4_K_M	4.9 GB	10 GB	★★★★★	Excellent multilingual — recommended
Phi-3-mini-4k-instruct Q4_K_M	2.2 GB	6 GB	★★★★	Lightweight machines — fast responses
Mixtral-8x7B-Instruct Q4_K_M	26 GB	48 GB	★★★★★	High precision — dedicated GPU server

All models are freely available on HuggingFace. Download directly from Admin → AI Model → Recommended models.

Licensing & Pricing

License	Price	Duration	Sessions	Users	Updates included	Support
Solo	€39	5-year	5	1	5 years	Email
Student	€19	5-year	5	1	5 years	Email

License	Price	Duration	Sessions	Users	Updates included	Support
Pro	€149/yr	Annual	50	5	Yes	Priority email
Business	€499/yr	Annual	Unlimited	Unlimited	Yes	Priority support

For multi-site environments, custom licensing is available. Contact: contact@fdevelopment.eu

Supported Formats & Collections

Format	Description	Collection	Typical use case
PDF	Contracts, reports, guides, case law, regulations	Internal contracts	Query the entire contract portfolio
DOCX	Word documents, internal procedures, notes	Compliance / GDPR	Policies, compliance procedures, DPAs
XLSX	Spreadsheets, tabular data, reference tables	Regulatory	Legislation, internal case law, standards
TXT / MD	Notes, technical documentation, internal wikis	HR Procedures	Employee handbook, job descriptions, onboarding

Security & GDPR Compliance

Zero network transmission

No data (documents, questions, responses) leaves your infrastructure. The AI model is a local file with no Internet access.

Privacy by Design

No telemetry, no tracking, no cloud subcontractors. Native GDPR compliance by architecture.

Encrypted conversations

Session history is automatically encrypted on disk. System administrators cannot read content in plain text.

Offline mode (air-gap)

Kno runs on isolated networks. No Internet connection required after initial installation.

Admin authentication

Admin interface protected by password. All active sessions immediately invalidated on password change.

No cloud subcontractor

No DPA (Data Processing Agreement) to sign. No transfer to third countries. Storage controlled by your organization.

Network Deployment

Mode	Description	Access	Configuration
Single workstation	Kno on the user's machine	localhost:8000	None (default)
Local server	Dedicated server, simultaneous multi-user access	http://192.168.x.x:8000	Windows Firewall / ufw Linux
HTTPS (optional)	nginx or Caddy reverse proxy in front of Kno	https://kno.domain.local	nginx / Caddy → localhost:8000
Custom port	If port 8000 is already in use	Any port	"port": 9000 in config.json

AI Engine Parameters (changeable without restart)

Parameter	Default	Range	Impact & recommendation
Retrieved excerpts (K)	3	1 – 10	Increase to 5–7 for dense documents (PDF tables). Balances completeness vs. speed.
Chunk size	700	200 – 2000	Reduce to 400–500 for tabular PDFs. Increase for long continuous text.

Parameter	Default	Range	Impact & recommendation
Overlap	100	0 – 500	Increase to 150–200 to avoid context breaks between sections.
Temperature	0.1	0.0 – 1.0	0.1 = factual responses (recommended for documentation). 0.5 = more natural.
GPU layers (n_gpu_layers)	0	0 – 100+	Set 35+ on CUDA-compatible NVIDIA GPU. Speed multiplied by 10 to 20.

Installation & License Activation

No software prerequisites

Kno is a standalone executable (.exe Windows / Linux binary). Python, Node.js, Docker are not required.

Installation steps

1. Download Kno-Setup-vX.X.X.exe from GitHub Releases
2. Run the installer (choose folder, shortcut)
3. Place a GGUF model in the models/ folder
4. Activate the license via Admin → License

Activation procedure

- 01 Launch Kno — the License interface opens automatically
- 02 Note your Machine ID (e.g.: B2C798229587EDE1)
- 03 Send this ID to contact@fdevelopment.eu from your PayPal email
- 04 Receive your activation code within 24 business hours
- 05 Paste the code and click Activate — no restart required

The license is tied to the machine ID. To change machines, contact support.